

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: house_sales=pd.read_csv(r"C:\Users\s323\Desktop\Gatherings\Data Science\ML\Amit Mi
```

```
In [3]: house_sales
```

```
Out[3]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0
...	...	...	...	...	...	...	...	...
21608	263000018	20140521T000000	360000.0	3	2.50	1530	1131	3.0
21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5813	2.0
21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1350	2.0
21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0
21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1076	2.0

21613 rows × 21 columns

```
In [4]: house_sales.head()
```

*# without head function it will show first five and last five transaction*

```
Out[4]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

```
In [5]: house_sales.shape
```

```
Out[5]: (21613, 21)
```

In [6]: `house_sales.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price               21613 non-null  float64
3   bedrooms            21613 non-null  int64
4   bathrooms           21613 non-null  float64
5   sqft_living         21613 non-null  int64
6   sqft_lot            21613 non-null  int64
7   floors              21613 non-null  float64
8   waterfront          21613 non-null  int64
9   view                21613 non-null  int64
10  condition            21613 non-null  int64
11  grade                21613 non-null  int64
12  sqft_above          21613 non-null  int64
13  sqft_basement       21613 non-null  int64
14  yr_built            21613 non-null  int64
15  yr_renovated        21613 non-null  int64
16  zipcode             21613 non-null  int64
17  lat                 21613 non-null  float64
18  long                21613 non-null  float64
19  sqft_living15       21613 non-null  int64
20  sqft_lot15          21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

In [7]: `house_sales.columns`

Out[7]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft\_living', 'sqft\_lot', 'floors', 'waterfront', 'view', 'condition', 'grade', 'sqft\_above', 'sqft\_basement', 'yr\_built', 'yr\_renovated', 'zipcode', 'lat', 'long', 'sqft\_living15', 'sqft\_lot15'], dtype='object')

## Data Wrangling

- Check if there is any missing value


In [8]: `house_sales.isnull().sum()`

```
Out[8]: id          0
        date        0
        price       0
        bedrooms    0
        bathrooms   0
        sqft_living  0
        sqft_lot     0
        floors      0
        waterfront  0
        view        0
        condition   0
        grade       0
        sqft_above   0
        sqft_basement 0
        yr_built     0
        yr_renovated 0
        zipcode      0
        lat          0
        long         0
        sqft_living15 0
        sqft_lot15   0
        dtype: int64
```


```
In [9]: # We have to predict house price, In linear regression - Assumption 1. Our target is price
        # Check the normality
```

## Assumptions

```
In [10]: from scipy.stats import norm
        sns.set_style("whitegrid")
        plt.figure(figsize=(13,6))
        sns.distplot(house_sales["price"], fit=norm,color="blue")
        plt.show()
```




```
In [11]: # Since the data is +vely skewed, need to exclude the outliers
        # need to check if there is any outliers
        plt.figure(figsize=(12,8))
        sns.set_style("whitegrid")
        sns.boxplot(house_sales["price"], color="blue")
        plt.show()
```



```
In [12]: # we can't remove all outliers since there are a lot and we will lose the data
```

```
In [13]: # QQ PLOT
```

```
from scipy import stats
plt.figure(figsize=(10,8))
sns.set_style("whitegrid")
stats.probplot(house_sales["price"], plot=plt)
plt.show()
```



## Handling outliers= we can do it by Z score

- We can do it by IQR as well, but it targets only one column at a time while Z-score is applied on whole data range

```
In [14]: house_sales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price                21613 non-null  float64
3   bedrooms             21613 non-null  int64
4   bathrooms            21613 non-null  float64
5   sqft_living          21613 non-null  int64
6   sqft_lot             21613 non-null  int64
7   floors               21613 non-null  float64
8   waterfront           21613 non-null  int64
9   view                 21613 non-null  int64
10  condition            21613 non-null  int64
11  grade                21613 non-null  int64
12  sqft_above           21613 non-null  int64
13  sqft_basement        21613 non-null  int64
14  yr_built             21613 non-null  int64
15  yr_renovated         21613 non-null  int64
16  zipcode              21613 non-null  int64
17  lat                  21613 non-null  float64
18  long                 21613 non-null  float64
19  sqft_living15        21613 non-null  int64
20  sqft_lot15          21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
In [15]: # Since there are 2 coluns which are not playing any major role in data, will drop
house_df=house_sales.drop(["id","date"],axis=1)
```

```
In [16]: house_df.head()
```

```
Out[16]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade
0	221900.0	3	1.00	1180	5650	1.0	0	0	3	
1	538000.0	3	2.25	2570	7242	2.0	0	0	3	
2	180000.0	2	1.00	770	10000	1.0	0	0	3	
3	604000.0	4	3.00	1960	5000	1.0	0	0	5	
4	510000.0	3	2.00	1680	8080	1.0	0	0	3	

```
In [17]: from scipy import stats
z=stats.zscore(house_df)

# for outliers z score should be greater than value - +3 or -3
```

```
In [18]: np.where(np.abs(z>4))
# np.where will return the row and col value
```

```
Out[18]: (array([ 1,  5, 21, ..., 21576, 21576, 21576], dtype=int64),
array([13, 17,  7, ...,  0,  6,  7], dtype=int64))
```

```
In [19]: # Let's see how much data we are loosing
len(np.where(np.abs(z>4))[0])
```

```
Out[19]: 2527
```


```
In [20]: # in percentage
len(np.where(np.abs(z>4))[0]) / len(house_df)
```

*# ideal case is 3, but here we are doing it at 4 but can't increase more than 4*

```
Out[20]: 0.11692037199833434
```

```
In [21]: house_df.drop(np.where(np.abs(z>4))[0], axis = 0, inplace=True)
```


```
In [22]: from scipy.stats import norm
sns.set_style("whitegrid")
plt.figure(figsize=(11,8))
sns.distplot(house_df["price"], fit=norm, color="crimson")
plt.show()
```



## Corelation

```
In [23]: plt.figure(figsize=(15,10))
sns.heatmap(house_df.corr(), annot=True, cmap="RdPu")
```

```
Out[23]: <AxesSubplot:>
```



In [24]: *# from the above diagram we can feel, two colomns isnot playing any role in correla*

In [25]: *# unique value to know if it consists any of the value*  
 house\_df['waterfront'].unique()


Out[25]: array([0], dtype=int64)

In [26]: house\_df['yr\_renovated'].unique()

Out[26]: array([0], dtype=int64)

In [27]: plt.figure(figsize=(3,10))  
 sns.heatmap(house\_df.corr()[["price"]],annot=True,cmap="RdPu")  
 plt.show()





In [28]: *# Any feature closer to 0, will be eliminated*

## Split X & Y

In [29]: `X = house_df.drop(["waterfront", "yr_renovated", "price"], axis=1)`  
*# Rest we will not drop weaker features now, because we want to do Lasso on it, it will*

In [30]: `Y = house_df["price"]`

## Train and Test Split

In [31]: `from sklearn.model_selection import train_test_split`  
*# choose a random state (0-11) 80% train , 20% test*  
*# random\_state = (1-11)*  
*# X is denoting feature and y is representing target*  
`x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=1)`

## Data Preprocessing, Data Scalar

```
In [32]: from sklearn.preprocessing import MinMaxScaler
# MinMaxScaler - scale all the features with range (0,1)
scaler=MinMaxScaler(feature_range=(0,1))
# fit + transform
x_train_scaler = scaler.fit_transform(x_train)
# only transform for test
x_test_scaler = scaler.transform(x_test)
```

## LASSO Regression

```
In [33]: from sklearn.linear_model import Lasso
# hyperparameters : alpha = (0.00001-1), max_iter = higher if sample size is small
lasso_model=Lasso(alpha = 0.0001, max_iter = 100000)

# Lasso model is created afterwards we do need to fit that
```

```
In [34]: lasso_model.fit(x_train_scaler, y_train)
```

```
Out[34]: Lasso(alpha=0.0001, max_iter=100000)
```

```
In [35]: # Algorithm is untrained and later we trained them and make them model to predict
lasso_model.score(x_test_scaler, y_test)
```

```
Out[35]: 0.7049161200498151
```

## Regression Metrics

- It is used for check the performance of regression algorithm ( also called accuracy matrix )
- MSE, MSA

```
In [36]: # making new Predictions (yhat)
yhat = lasso_model.predict(x_test_scaler)
```

```
In [37]: # r2 SCORE, MSE, MAE
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
# coefficient of determination (0 - 1)
r2_score(y_test, yhat)
```

```
Out[37]: 0.7049161200498151
```

```
In [38]: mean_absolute_error(y_test, yhat)
```

```
Out[38]: 102494.56640321515
```

```
In [39]: # RMSE
np.sqrt(mean_squared_error(y_test, yhat))
```

*# here we can see that the RMSE value is way bigger but still we need to see actual*

```
Out[39]: 147902.42875551208
```

```
In [40]: Y.mean()
# quite away from original mean
```

```
Out[40]: 499567.6842822532
```

## Feature selection using Lasso

- how does lasso useful in feature selection

```
In [41]: # lasso.coef_ is beta value such as
# Lasso model will shrink the Least imp't feature's coefficient (b0, b1,b2) closer to 0
lasso_model.coef_
```

```
Out[41]: array([-134752.6994601 , 159438.26741008, 1101684.57475157,
        36468.41885377, 69724.91743785, 130568.19229457,
        111048.62667001, 994031.93381119, -435980.42470607,
        -224314.29930833, -269815.46292146, -79402.67035199,
        346568.24409264, -91777.0024419 , 166977.06157302,
        -104014.7823389 ])
```


```
In [42]: # in order to know that we had to create a dataframe, of columns and columns coeff
# creating empty data frame and then adding columns with values
lasso_coef = pd.DataFrame()
lasso_coef['Columns'] = x_train.columns
lasso_coef['Coefficient Estimate'] = pd.Series(lasso_model.coef_)
print(lasso_coef)
```

```
# b0 won't be there since it is for intializing
```

	Columns	Coefficient Estimate
0	bedrooms	-1.347527e+05
1	bathrooms	1.594383e+05
2	sqft_living	1.101685e+06
3	sqft_lot	3.646842e+04
4	floors	6.972492e+04
5	view	1.305682e+05
6	condition	1.110486e+05
7	grade	9.940319e+05
8	sqft_above	-4.359804e+05
9	sqft_basement	-2.243143e+05
10	yr_built	-2.698155e+05
11	zipcode	-7.940267e+04
12	lat	3.465682e+05
13	long	-9.177700e+04
14	sqft_living15	1.669771e+05
15	sqft_lot15	-1.040148e+05

```
In [43]: # columns vs their coeff and closer to 0, we will eliminate them
# we can play with alpha value
plt.figure(figsize=(20,9))
plt.bar(lasso_coef['Columns'],lasso_coef['Coefficient Estimate'])
```

```
Out[43]: <BarContainer object of 16 artists>
```



```
In [44]: # closer to 0, squarefit plot and zipcode can be eliminated
# we can change with alpha value as well
# alpha is learning rate
```

## Lasso Regression with alpha = 1

- Basically we will retrain our model with lasso at alpha = 1

```
In [45]: from sklearn.linear_model import Lasso
# hyperparameters : alpha = (0.00001-1), max_iter = higher if sample size is small
lasso_model=Lasso(alpha = 1, max_iter = 100000)
```

```
In [46]: lasso_model.fit(x_train_scaler, y_train)
```

```
Out[46]: Lasso(alpha=1, max_iter=100000)
```

```
In [47]: # Algorithm is untrained and later we trained them and make them model to predict
lasso_model.score(x_test_scaler, y_test)
```


```
Out[47]: 0.7049159098093531
```

```
In [48]: lasso_coef = pd.DataFrame()
lasso_coef['Columns'] = x_train.columns
lasso_coef['Coefficient Estimate'] = pd.Series(lasso_model.coef_)
print(lasso_coef)
```

	Columns	Coefficient Estimate
0	bedrooms	-134585.475691
1	bathrooms	159338.157479
2	sqft_living	510405.741327
3	sqft_lot	35770.376987
4	floors	69717.404177
5	view	130561.199316
6	condition	111009.768583
7	grade	993950.584686
8	sqft_above	82304.299707
9	sqft_basement	0.000000
10	yr_built	-269764.784051
11	zipcode	-79363.261753
12	lat	346558.381581
13	long	-91721.988286
14	sqft_living15	166949.531808
15	sqft_lot15	-103392.745685

```
In [49]: plt.figure(figsize=(20,9))
plt.bar(lasso_coef['Columns'],lasso_coef['Coefficient Estimate'])
```

Out[49]: <BarContainer object of 16 artists>



```
In [50]: # waterfront, yearrenovated, price are from starting and remaining 2 features from
X = house_df.drop(['waterfront', 'sqft_basement', 'sqft_lot15', 'yr_renovated', 'price'])
Y = house_df['price']
```

```
In [51]: x_train,x_test,y_train,y_test=train_test_split(X,Y, test_size=0.2, random_state=8)
```

```
In [52]: scaler=MinMaxScaler(feature_range=(0,1))
#fit + transform
x_train.scaler=scaler.fit_transform(x_train)
# only transform for test
x_test.scaler=scaler.transform(x_test)
```

```
In [53]: lasso_model = Lasso(alpha = 1, max_iter = 100000)
```

```
In [54]: lasso_model.fit(x_train.scaler, y_train)
```

Out[54]: Lasso(alpha=1, max\_iter=100000)

```
In [55]: lasso_model.score(x_test.scaler, y_test)
```

Out[55]: 0.7049159098093531

```
In [56]: # We can observe very minor change is with dropping the feature and without dropping
```

## Ridge Regression

```
In [57]: from sklearn.linear_model import Ridge
# hyperparameters : alpha = (0.00001-1), max_iter = higher if sample size is small
ridge_model = Ridge(alpha = 0.01, max_iter = 100000)
```

```
In [58]: ridge_model.fit(x_train.scaler,y_train)
```

Out[58]: Ridge(alpha=0.01, max\_iter=100000)

```
In [59]: ridge_model.score(x_test.scaler,y_test)
```

Out[59]: 0.7044084249552689

## Ridge Regression Metrics

```
In [60]: # making new Predictions (yhat)
#yhat = ridge_model.predict(x_test_scaler)

#yhat = ridge_model.predict(x_test_scaler)
```

```
In [61]: # r2-score , mean_squared_error, mean_absolute_error
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
# coefficient of determination (0 - 1)
r2_score(y_test, yhat)
```

Out[61]: 0.7049161200498151

```
In [62]: mean_absolute_error(y_test, yhat)
```

Out[62]: 102494.56640321515

```
In [63]: # RMSE
np.sqrt(mean_squared_error(y_test, yhat))
```

Out[63]: 147902.42875551208

## Elastic net- Combination of both Lasso and Ridge model

```
In [64]: from sklearn.linear_model import ElasticNet
elasticNet_model = ElasticNet(alpha = 0.004, max_iter=100000)
```

```
In [65]: elasticNet_model.fit(x_train_scaler, y_train)
```

Out[65]: ElasticNet(alpha=0.004, max\_iter=100000)

```
In [66]: elasticNet_model.score(x_test_scaler, y_test)
```

Out[66]: 0.6993882344502862

In [ ]: